

sns-ep-1-2-3-aaronmedinamelian

January 28, 2024

```
[53]: import pandas as pd
from sklearn.feature_selection import SelectKBest, f_classif

# Lee el archivo CSV
archivo_csv = "obesity.csv"
datos = pd.read_csv(archivo_csv)

datos
```

```
[53]:
```

	Gender	Age	family_history_with_overweight	FAVC	FCVC	NCP	\
0	0	21.000000	1	0	2.0	3.0	
1	0	21.000000	1	0	3.0	3.0	
2	1	23.000000	1	0	2.0	3.0	
3	1	27.000000	0	0	3.0	3.0	
4	1	22.000000	0	0	2.0	1.0	
...
2106	0	20.976842	1	1	3.0	3.0	
2107	0	21.982942	1	1	3.0	3.0	
2108	0	22.524036	1	1	3.0	3.0	
2109	0	24.361936	1	1	3.0	3.0	
2110	0	23.664709	1	1	3.0	3.0	

	CAEC	SMOKE	CH2O	SCC	FAF	TUE	CALC	Automobile	Bike	\
0	1	0	2.000000	0	0.000000	1.000000	0	0	0	
1	1	1	3.000000	1	3.000000	0.000000	1	0	0	
2	1	0	2.000000	0	2.000000	1.000000	2	0	0	
3	1	0	2.000000	0	2.000000	0.000000	2	0	0	
4	1	0	2.000000	0	0.000000	0.000000	1	0	0	
...
2106	1	0	1.728139	0	1.676269	0.906247	1	0	0	
2107	1	0	2.005130	0	1.341390	0.599270	1	0	0	
2108	1	0	2.054193	0	1.414209	0.646288	1	0	0	
2109	1	0	2.852339	0	1.139107	0.586035	1	0	0	
2110	1	0	2.863513	0	1.026452	0.714137	1	0	0	

	Motorbike	Public_Transportation	Walking	NObeyesdad
0	0	1	0	0

1	0	1	0	0
2	0	1	0	0
3	0	0	1	0
4	0	1	0	0
...
2106	0	1	0	1
2107	0	1	0	1
2108	0	1	0	1
2109	0	1	0	1
2110	0	1	0	1

[2111 rows x 19 columns]

```
[54]: # Itera sobre todas las columnas y muestra los valores y sus frecuencias
for columna in datos.columns:
    conteo_valores = datos[columna].value_counts()
    print(f"\nColumna: {columna}\n{conteo_valores}")
```

Columna: Gender

Gender

1 1068

0 1043

Name: count, dtype: int64

Columna: Age

Age

18.000000 128

26.000000 101

21.000000 96

23.000000 89

19.000000 59

...

23.320120 1

34.243146 1

18.549437 1

36.310292 1

23.664709 1

Name: count, Length: 1402, dtype: int64

Columna: family_history_with_overweight

family_history_with_overweight

1 1726

0 385

Name: count, dtype: int64

Columna: FAVC

FAVC

1 1866

0 245

Name: count, dtype: int64

Columnna: FCVC

FCVC

3.000000 652

2.000000 600

1.000000 33

2.823179 2

2.214980 2

...

2.927409 1

2.706134 1

2.010684 1

2.300408 1

2.680375 1

Name: count, Length: 810, dtype: int64

Columnna: NCP

NCP

3.000000 1203

1.000000 199

4.000000 69

2.776840 2

3.985442 2

...

3.054899 1

3.118013 1

3.335876 1

3.205009 1

1.089048 1

Name: count, Length: 635, dtype: int64

Columnna: CAEC

CAEC

1 1765

2 242

3 53

0 51

Name: count, dtype: int64

Columnna: SMOKE

SMOKE

0 2067

1 44

Name: count, dtype: int64

Columna: CH20

CH20

2.000000	448
1.000000	211
3.000000	162
2.825629	3
1.636326	3

...

1.622638	1
2.452986	1
2.035954	1
1.944095	1
2.863513	1

Name: count, Length: 1268, dtype: int64

Columna: SCC

SCC

0	2015
1	96

Name: count, dtype: int64

Columna: FAF

FAF

0.000000	411
1.000000	234
2.000000	183
3.000000	75
0.110174	2

...

1.916751	1
0.954459	1
0.340915	1
0.986414	1
1.026452	1

Name: count, Length: 1190, dtype: int64

Columna: TUE

TUE

0.000000	557
1.000000	292
2.000000	109
0.630866	4
1.119877	3

...

1.343044	1
1.019452	1
0.673408	1

```
0.997600      1
0.714137      1
Name: count, Length: 1129, dtype: int64
```

```
Columnna: CALC
CALC
1      1401
0       639
2        70
3         1
Name: count, dtype: int64
```

```
Columnna: Automobile
Automobile
0      1654
1       457
Name: count, dtype: int64
```

```
Columnna: Bike
Bike
0      2104
1         7
Name: count, dtype: int64
```

```
Columnna: Motorbike
Motorbike
0      2100
1        11
Name: count, dtype: int64
```

```
Columnna: Public_Transportation
Public_Transportation
1      1580
0       531
Name: count, dtype: int64
```

```
Columnna: Walking
Walking
0      2055
1         56
Name: count, dtype: int64
```

```
Columnna: NObeyesdad
NObeyesdad
0      1139
1       972
Name: count, dtype: int64
```

- FCVC realmente se encuentra en una escala de 1 a 3
- NCP realmente se encuentra en una escala de 1 a 4

Por lo tanto, he decidido redondear estos valores para asemejarlos lo más posible a la documentación proporcionada.

```
[55]: # Nombre de las columnas a redondear
columnas_a_redondear = ["FCVC", "NCP"]

# Redondear los valores de las columnas específicas a enteros
datos[columnas_a_redondear] = datos[columnas_a_redondear].round().astype(int)
```

```
[56]: # Itera sobre todas las columnas y muestra los valores y sus frecuencias
for columna in datos.columns:
    conteo_valores = datos[columna].value_counts()
    print(f"\nColumna: {columna}\n{conteo_valores}")
```

Columna: Gender

Gender

1 1068

0 1043

Name: count, dtype: int64

Columna: Age

Age

18.000000 128

26.000000 101

21.000000 96

23.000000 89

19.000000 59

...

23.320120 1

34.243146 1

18.549437 1

36.310292 1

23.664709 1

Name: count, Length: 1402, dtype: int64

Columna: family_history_with_overweight

family_history_with_overweight

1 1726

0 385

Name: count, dtype: int64

Columna: FAVC

FAVC

1 1866

0 245
Name: count, dtype: int64

Columnna: FCVC

FCVC

2 1013

3 996

1 102

Name: count, dtype: int64

Columnna: NCP

NCP

3 1470

1 316

2 176

4 149

Name: count, dtype: int64

Columnna: CAEC

CAEC

1 1765

2 242

3 53

0 51

Name: count, dtype: int64

Columnna: SMOKE

SMOKE

0 2067

1 44

Name: count, dtype: int64

Columnna: CH2O

CH2O

2.000000 448

1.000000 211

3.000000 162

2.825629 3

1.636326 3

...

1.622638 1

2.452986 1

2.035954 1

1.944095 1

2.863513 1

Name: count, Length: 1268, dtype: int64

Columnna: SCC

```
SCC
0    2015
1      96
Name: count, dtype: int64
```

```
Columnna: FAF
FAF
0.000000    411
1.000000    234
2.000000    183
3.000000     75
0.110174      2
...
1.916751      1
0.954459      1
0.340915      1
0.986414      1
1.026452      1
Name: count, Length: 1190, dtype: int64
```

```
Columnna: TUE
TUE
0.000000    557
1.000000    292
2.000000    109
0.630866      4
1.119877      3
...
1.343044      1
1.019452      1
0.673408      1
0.997600      1
0.714137      1
Name: count, Length: 1129, dtype: int64
```

```
Columnna: CALC
CALC
1    1401
0     639
2      70
3       1
Name: count, dtype: int64
```

```
Columnna: Automobile
Automobile
0    1654
1     457
Name: count, dtype: int64
```



```
Columna: Bike
Bike
0    2104
1         7
Name: count, dtype: int64
```

```
Columna: Motorbike
Motorbike
0    2100
1         11
Name: count, dtype: int64
```

```
Columna: Public_Transportation
Public_Transportation
1    1580
0     531
Name: count, dtype: int64
```

```
Columna: Walking
Walking
0    2055
1         56
Name: count, dtype: int64
```

```
Columna: NObeyesdad
NObeyesdad
0    1139
1     972
Name: count, dtype: int64
```

```
[57]: columnas_seleccionadas = ['Age', 'family_history_with_overweight', 'FAVC', 'CAEC', 'SCC']

columnas_seleccionadas_con_obesidad = ['Age', 'family_history_with_overweight', 'FAVC', 'CAEC', 'SCC', 'NObeyesdad']
```

1 Examen/Proyecto UT1, UT2 y UT3

```
[59]: from sklearn.model_selection import train_test_split

train, test = train_test_split(datos[['Age', 'family_history_with_overweight', 'FAVC', 'CAEC', 'SCC', 'NObeyesdad']], test_size=0.33)

train.reset_index(inplace = True)
```

```
train
```

```
[59]:
```

	index	Age	family_history_with_overweight	FAVC	CAEC	SCC	\
0	1357	18.000000	1	1	1	0	
1	1971	19.297004	1	1	1	0	
2	1672	30.796262	1	1	1	0	
3	326	18.000000	1	1	1	0	
4	1118	18.198322	1	0	1	0	
...
1409	1369	18.078256	1	1	1	0	
1410	1025	38.464538	1	1	1	0	
1411	1618	41.000000	1	1	1	0	
1412	1937	21.140165	1	1	1	0	
1413	1178	19.684891	1	0	1	0	

```
NObayesdad
```

0	1
1	1
2	1
3	0
4	0
...	...
1409	1
1410	0
1411	1
1412	1
1413	0

```
[1414 rows x 7 columns]
```

- 1) (10%) Entrenar un modelo utilizando KNN. Realizar una búsqueda de los parámetros que ofrecen mejores resultados aplicando Cross Validation. Se valorará si se utiliza una gráfica para mostrar los diferentes niveles precisión para cada uno de los parámetros probados.

```
[60]: from sklearn import neighbors
from sklearn.model_selection import KFold
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
import numpy as np

# Para el entrenamiento se aplica Cross Validation (Validación cruzada)
cv = KFold(n_splits = 5, shuffle = True) # shuffle = False si hay dimensión
    ↪ temporal

for i, weights in enumerate(['uniform', 'distance']):
```

```

total_scores = []
for n_neighbors in range(1,50):
    fold_accuracy = []
    knn = neighbors.KNeighborsClassifier(n_neighbors, weights=weights) # La
    ↪ métrica por defecto es minkowski
    #knn = neighbors.KNeighborsClassifier(n_neighbors, weights=weights,
    ↪ metric="euclidean")
    for train_fold, test_fold in cv.split(train):
        # División train test aleatoria
        f_train = train.loc[train_fold]
        f_test = train.loc[test_fold]
        # entrenamiento y ejecución del modelo
        knn.fit( X = f_train.drop(['NObeyesdad'], axis=1),
                y = f_train['NObeyesdad'])
        y_pred = knn.predict(X = f_test.drop(['NObeyesdad'], axis = 1))
        # evaluación del modelo
        acc = accuracy_score(f_test['NObeyesdad'], y_pred)
        fold_accuracy.append(acc)
    total_scores.append(sum(fold_accuracy)/len(fold_accuracy))


plt.plot(range(1,len(total_scores)+1), total_scores,
         marker='o', label=weights)
print ('Max Value ' + weights + " : " + str(max(total_scores)) + " (" +
    ↪ str(np.argmax(total_scores) + 1) + ")")
plt.ylabel('Acc')

plt.legend()
plt.show()

```

Max Value uniform : 0.9731273337844272 (5)

Max Value distance : 0.9724256321579832 (6)



```
[61]: total_scores
```

```
[61]: [0.9632283286970905,
0.9603764128010426,
0.9681753251635214,
0.9667518732927348,
0.9674661053053655,
0.9724256321579832,
0.9681753251635214,
0.9709996742099591,
0.9717189183770645,
0.9703029847379897,
0.9696037891887828,
0.9724206200035086,
0.9710046863644337,
0.9717164122998272,
0.9724131017717965,
0.9702979725835149,
0.9702979725835149,
0.9695837405708844,
0.9710071924416711,
0.9710021802871964,
```

```

0.9695937648798336,
0.9710071924416711,
0.9702954665062776,
0.9688845450216774,
0.9702954665062776,
0.9709971681327219,
0.9702929604290403,
0.9709921559782473,
0.9702954665062776,
0.9681703130090469,
0.9703029847379897,
0.9695812344936471,
0.9717114001453524,
0.9709996742099591,
0.9688895571761522,
0.9695862466481217,
0.9695837405708844,
0.9702929604290403,
0.970305490815227,
0.9702954665062776,
0.9703004786607524,
0.9703004786607524,
0.9688895571761522,
0.9695912588025962,
0.9695987770343082,
0.9703004786607524,
0.9702979725835149,
0.9702979725835149,
0.9695862466481218]

```

```

[62]: # constructor
n_neighbors = 12
weights = 'distance'
knn = neighbors.KNeighborsClassifier(n_neighbors= n_neighbors, weights=weights)
    ↪ # Se utiliza este método porque es un problema de clasificación
# fit and predict
knn.fit(X = train[['Age', 'family_history_with_overweight', 'FAVC', 'CAEC',
    ↪ 'SCC']], y = train['NObeyesdad']) # En este caso no entrena sino memoriza
    ↪ los valores
y_pred = knn.predict(X = test[['Age', 'family_history_with_overweight', 'FAVC',
    ↪ 'CAEC', 'SCC']])
acc = accuracy_score(test['NObeyesdad'], y_pred)
print ('Acc', acc)

```

Acc 0.7489239598278336

```
[63]: from sklearn.metrics import classification_report
print(classification_report(test['NObeyesdad'], y_pred))
```

	precision	recall	f1-score	support
0	0.77	0.76	0.77	375
1	0.73	0.73	0.73	322
accuracy			0.75	697
macro avg	0.75	0.75	0.75	697
weighted avg	0.75	0.75	0.75	697

- 2) (10%) Entrenar un modelo utilizando alguna de las técnicas de árboles que hemos visto en clase. Si no ha dado tiempo de verlo en clase se propone realizar una pequeña investigación sobre posibles modelos y proponer uno que consideren adecuado para el entrenamiento. Justificando la elección del mismo.

```
[64]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report

X_tree = datos[columnas_seleccionadas_con_obesidad].drop('NObeyesdad', axis=1)
    ↪ # Características
y_tree = datos['NObeyesdad'] # Etiquetas de clase
```

```
[65]: X_train_tree, X_test_tree, y_train_tree, y_test_tree = train_test_split(X_tree,
    ↪ y_tree, test_size=0.2, random_state=42)

# Crear el modelo
model_tree = DecisionTreeClassifier(random_state=42)

# Entrenar el modelo
model_tree = model_tree.fit(X_train_tree, y_train_tree)
```

```
[67]: # Realizar predicciones en el conjunto de prueba
y_pred_tree = model_tree.predict(X_test_tree)

# Evaluar la precisión del modelo
accuracy_tree = accuracy_score(y_test_tree, y_pred_tree)
print(f'Precisión del modelo: {accuracy_tree:.2f}')

# Mostrar informe de clasificación
print('Informe de clasificación:')
print(classification_report(y_test_tree, y_pred_tree))
```

Precisión del modelo: 0.77

Informe de clasificación:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.75	0.86	0.80	224
1	0.81	0.67	0.74	199
accuracy			0.77	423
macro avg	0.78	0.77	0.77	423
weighted avg	0.78	0.77	0.77	423

- 3) (10%) De las tres técnicas utilizadas (NAIVEBAYES, KNN, ARBOLES), selecciona aquella que te ofrece mejores resultados de predicción. Argumenta tu elección.
- 4) (5%) Exporta el fichero .pkl que contiene el modelo y los parámetros utilizados durante el entrenamiento.

```
[68]: import sklearn.externals
import joblib
import pickle

joblib.dump(knn, 'KNN_model_obesity.pkl')
```

```
[68]: ['KNN_model_obesity.pkl']
```

- 5) (10%) Importa el modelo .pkl y realiza una predicción utilizando un conjunto de valores, definido por ti, asociados a las características con las que entrenaste el modelo.

```
[69]: # Para importar el modelo entrenado y ejecutar de nuevo test
clf_entrenado = neighbors.KNeighborsClassifier(50, weights="distance")
clf_entrenado = joblib.load('KNN_model_obesity.pkl')
```

```
[74]: clf_entrenado.predict(np.array([[40,1,0,0,0]]))
```

```
c:\Users\Beatriz\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X
does not have valid feature names, but KNeighborsClassifier was fitted with
feature names
  warnings.warn(
```

```
[74]: array([1], dtype=int64)
```